

# CSC Latin America 2026

---



## Tools & Techniques + Software Design

High-Performance Computing for Experimental Physics and Data Analysis

Arturo Sánchez Pineda – INAIT AI

January 2026

<https://indico.cern.ch/e/CSC2026-Chile>

# Welcome & Overview

---

## What you'll learn:

- **Tools & Techniques:** Correctness, collaboration, and performance fundamentals
- **Software Design:** Multi-core computing, data layout, and optimisation strategies
- **Practical Skills:** Hands-on exercises with real HEP workflows

## Why this matters:

Modern HEP experiments generate petabytes of data. Writing correct, maintainable, and performant code isn't optional — it's essential.

**Philosophy: Correctness first, then performance — always with collaboration in mind.**

# About Me



---

Arturo Sánchez Pineda – from Venezuela

- Senior DevOps Engineer and Infrastructure Responsible at INAIT AI, Switzerland.
- PhD in Physics, University of Naples Federico II.
- Former co-creator and coordinator of the ATLAS Open Data project at CERN.
- Former researcher/educator at ULA, ICTP, INFN, University of Udine.
- Co-founder of LA-CoNGA physics (Latin American alliance for Capacity-building in Advanced Physics), CEVALE2VE, and the Creative Commons Venezuela and Switzerland chapters.

**Mission:** Making scientific education accessible worldwide – especially in Latin America.

# Schedule at a Glance

---

Day	Session	Topic
Mon, Jan 12	12:30	T&T L1: Correctness first
	14:30	T&T L2: From my code to our code
	17:00–18:00	T&T E1–E2: Exercises
Tue, Jan 13	11:30	SD L1: Many-core foundations
Wed, Jan 14	09:45	SD L2: Patterns & data layout
	12:30	T&T L3: Performance fundamentals
	14:30–15:30	SD E1–E2: Exercises
Thu, Jan 15	11:30	SD L3: Profiling to optimisation

# Courses Philosophy

---

## Three Pillars of Quality Software:

1. **Correctness** — Does it produce the right answer?
2. **Maintainability** — Can others (including future you) work with it?
3. **Performance** — Does it run fast enough for production?

**In that order.** A fast, wrong answer is useless. An unmaintainable codebase dies.

**My approach:** Learn tools and techniques that enforce correctness, enable collaboration, and unlock performance — systematically.

**Fun fact:** This course uses real **ATLAS Open Data at 13 TeV**. (**The Higgs boson was discovered with 7-8 TeV data recorded in 2011-2012.**)

# Learning Objectives

---

By the end of this course, you will:

- Debug efficiently with sanitisers (ASan/UBSan/TSan)
- Write robust tests with Catch2/pytest
- Collaborate at scale with Git workflows and CI/CD
- Measure performance scientifically
- Parallelise code correctly with OpenMP/TBB
- Optimise systematically
- Work with real HEP data (ROOT + ATLAS Open Data)

**Assessment:** Hands-on exercises. You'll **do** these things, not just hear about them.

# Tools We'll Use

---

## Languages & Frameworks:

- C++17/20, Python 3.x
- ROOT (CERN's data analysis framework)
- Catch2 (C++ testing), pytest (Python testing)

## Build & CI:

- CMake, Git, GitHub Actions, GitLab CI

## Performance & Debugging:

- GDB, ASan/UBSan/TSan, perf, Google Benchmark
- OpenMP, oneTBB (Threading Building Blocks)

## Infrastructure Context:

- CVMFS, EOS, Apptainer (Singularity)

# TOOLS & TECHNIQUES

---

## Lecture 1: Correctness First

"Make it work, make it right, make it fast — in that order"

60 minutes

Learn to catch bugs before they become disasters

# T&T L1: Correctness First

---

## Why Correctness Matters

### Real ATLAS examples:

**Case 1:** A memory leak in track reconstruction running on the Worldwide LHC Computing Grid (WLCG) – 500,000 cores processing data 24/7. One leak wastes **years of CPU time overnight.**

**Case 2:** An off-by-one error in event selection code. Months of analysis were invalidated. Paper submission delayed. Career impact for PhD students.

**Case 3:** Race condition in parallel Higgs analysis. Intermittent wrong results that only appear in 1% of runs. Nearly impossible to debug without proper tools.

**In HEP:** We're not just writing code — we're writing **scientific instruments**. Bugs don't just crash programs; they corrupt physics results and waste years of work.

# The Testing Mindset

---

**Interactive question:** *Raise your hand if you've ever...*

- Spent hours debugging, then found the bug was a typo?
- Had code that "worked on your machine" but failed elsewhere?
- Refactored code and accidentally broke something that was working?
- Trusted a result, then found it was wrong months later?

**The pattern:** All these are preventable with testing.

# The Testing Mindset

---

**Question:** When should you write tests?

**Wrong Answers:**

- "After the code works"
- "When I have time"
- "For important code only"
- "When someone asks"

**Right Answer:**

- **Always**
- **From the start**
- **For all code**
- **Because science**

**Test-Driven Development (TDD):** Write the test first, watch it fail, make it pass, refactor.

# Testing Frameworks: Catch2 (C++)

```
#include <catch2/catch_test_macros.hpp>

TEST_CASE("Energy is correct", "[physics]") {
    Particle p{1.0, 2.0, 3.0, 0.139};
    REQUIRE(p.energy() = Catch::Approx(3.744).epsilon(0.001));
}

TEST_CASE("Momentum conserved", "[physics]") {
    Particle p1{1.0, 0.0, 0.0, 0.139};
    Particle p2{-1.0, 0.0, 0.0, 0.139};
    REQUIRE((p1 + p2).px() = Catch::Approx(0.0).margin(1e-5));
}
```

Run: cmake --build build && ctest

# Testing Frameworks: pytest (Python)

```
from analysis import ParticleSelector

def test_pt_cut():
    sel = ParticleSelector(pt_min=20.0)
    events = [{"pt": 15.0}, {"pt": 25.0}, {"pt": 30.0}]

    result = sel.apply(events)
    assert len(result) == 2
    assert all(e["pt"] >= 20.0 for e in result)

def test_empty_input():
    sel = ParticleSelector(pt_min=20.0)
    assert sel.apply([]) == []
```

Run: `pytest -v`

# The Holy Trinity of Sanitisers

---

## Three Tools That Change Everything

### ASan (AddressSanitizer)

- Detects memory errors
- Buffer overflows
- Use-after-free
- Memory leaks

### TSan (ThreadSanitizer)

- Finds data races
- Detects race conditions
- Thread safety violations
- Lock order inversions

### UBSan (UndefinedBehaviorSanitizer)

- Catches undefined behaviour
- Integer overflow
- Null pointer dereference
- Division by zero

# Sanitisers: A Demonstration

---

## Without ASan (Traditional Debugging)

```
$ gdb ./reconstruction
(gdb) run
Program received signal SIGSEGV, Segmentation fault.
0x00007ffff7a5d123 in ?? ()
(gdb) bt
#0  0x00007ffff7a5d123 in ?? ()
#1  0x0000000000000000 in ?? ()
```

Time to find bug: 15–45 minutes of detective work

# Sanitisers: A Demonstration

## With ASan (Modern Debugging)

```
$ ./reconstruction
==12345==ERROR: AddressSanitizer: heap-buffer-overflow on address
0x603000000044 at pc 0x0000004a2b76 bp 0x7fffffff950 sp 0x7fffffff948
READ of size 4 at 0x603000000044 thread T0
#0 0x4a2b75 in TrackReconstructor::processHit(Hit const&)
    reconstruction.cpp:142
#1 0x4a1f32 in TrackReconstructor::run()
    reconstruction.cpp:87
#2 0x4a0d21 in main main.cpp:23

0x603000000044 is located 4 bytes to the right of 64-byte region
allocated by thread T0 here:
#0 0x4a3e88 in operator new(unsigned long)
```

Time to find bug: Instant, with exact line number

# Demo: Sanitisers in Action

---

## The Bug

```
void processEvents(const std::vector<Event>& events) {  
    float energies[100];  
  
    for (size_t i = 0; i ≤ events.size(); ++i) { // BUG: ≤ should be <  
        energies[i] = events[i].totalEnergy();  
    }  
    // More processing ...  
}
```

### What happens:

- Without sanitisers: Random crashes, corrupted data, or "works on my machine"
- With ASan: Immediate error message pointing to line with `≤`

# Common Debugging Horror Stories

---

## The Heisenbug:

"Crashes only on Fridays after 3 pm when I'm not watching."

→ Race condition. TSan catches it instantly.

## "It Worked Yesterday":

"I didn't change anything!"

→ Uninitialised memory. UBSan catches it.

## The Performance Cliff:

"Gets slower and slower, then crashes after 6 hours."

→ Memory leak. ASan catches it.

**Lesson:** Sanitisers turn mysterious bugs into obvious errors.

# Static Analysis: Catching Bugs Before Runtime

## clang-tidy

**What it does:** Analyses code without running it, finds potential bugs and style violations

```
// Code:  
void updateMomentum(Particle* p) {  
    if (p->mass() = 0.0) { // BUG: = instead of ==  
        p->setMomentum(0, 0, 0);  
    }  
}
```

clang-tidy output:

```
warning: using the result of an assignment as a condition  
        without parentheses [bugprone-assignment-in-if-condition]  
if (p->mass() = 0.0) {  
~~~~~^~~~~~
```

# C++ Core Guidelines

---

## Modern C++ Best Practices

Key principles we'll follow:

- **P.1:** Express ideas directly in code
- **P.5:** Prefer compile-time checking to run-time checking
- **F.7:** For general use, prefer `unique_ptr` to raw pointers
- **ES.20:** Always initialise an object
- **C.45:** Don't define a default constructor that only initialises data members

Tool: `clang-tidy` with `-checks=cppcoreguidelines-*`

Resources: <https://isocpp.github.io/CppCoreGuidelines/>

# Reproducible Environments

---

## The "Works On My Machine" Killer

**Common scenario:**

- Developer: "The code works perfectly on my laptop"
- CI:  Build failed
- Colleague: "I get segfaults when I run it"
- Production: Catastrophic failure!

**Root cause:** Environment differences (compiler versions, libraries, OS quirks)

**Solution:** Dev Containers + Docker

# Dev Containers

---

## Containerised Development Environment

**.devcontainer/devcontainer.json :**

```
{  
  "name": "HEP Analysis Environment",  
  "image": "gitlab-registry.cern.ch/linuxsupport/alma9-base:latest",  
  "features": {  
    "ghcr.io/devcontainers/features/git:1": {},  
    "ghcr.io/devcontainers/features/common-utils:2": {}  
  },  
  "postCreateCommand": "bash -lc 'source /cvmfs/sft.cern.ch/lcg/views/LCG_104/x86_64-el9-gcc13-opt/setup.sh'",  
  "customizations": {  
    "vscode": {  
      "extensions": ["ms-vscode.cpptools", "ms-python.python"]  
    }  
  }  
}
```

# HEP Context: CVMFS vs Containers

---

## CVMFS (CernVM-FS)

- **Purpose:** Software distribution
- Read-only file system
- Cached globally via Squid
- `/cvmfs/sft.cern.ch/`
- Perfect for: Sharing LCG releases

## Containers (Apptainer/Docker)

- **Purpose:** Runtime isolation
- User namespaces
- Portable execution
- Bundled dependencies
- Perfect for: Reproducible jobs

**Best practice:** CVMFS for distribution layer + Containers for runtime layer

# Summary: T&T L1

---

## Key Takeaways:

1. **Test everything** – Use Catch2/pytest from day one
2. **Sanitisers save time** – ASan/UBSan/TSan find bugs instantly
3. **Static analysis catches bugs early** – Run clang-tidy in CI
4. **Reproducibility matters** – Dev Containers eliminate environment issues
5. **HEP infrastructure** – CVMFS for distribution, containers for execution

**Next:** How do we move from "my code" to "our code"?

# TOOLS & TECHNIQUES

---

## Lecture 2: From My Code to Our Code

"Individual brilliance doesn't scale — collaboration does"

60 minutes

Master Git, CI/CD, and workflows for 3,000-person teams

# T&T L2: From My Code to Our Code

---

## Motivation

Individual brilliance doesn't scale.

In HEP experiments:

- 3,000+ collaborators across 40+ countries
- Millions of lines of code (Athena, CMSSW)
- Code from 2008 is still running!
- 6 continents, 12 time zones

**Challenge:** Write code that survives collaboration

**Example:** ATLAS Athena: 6M lines of C++/Python, 500+ developers, 20k commits/year.

# Git for Large Collaborations

---

## Branching Models

### Feature Branch Workflow:

```
main
└── feature/track-reco
└── feature/electron-id
└── hotfix/memory-leak
```

Used by: CMSSW, Athena

### Git Flow:

```
main (production)
└── develop
    └── feature/xyz
        └── release/v1.2
```

Used by: Many ROOT projects

**Key principle:** `main` is always deployable. Development happens in branches.

# Commit Hygiene

---

## Atomic Commits

### Bad commit:

```
commit abc123
"Fixed stuff"
- Fixed track reconstruction bug
- Updated documentation
- Refactored particle class
- Added new analysis
- Cleaned up whitespace
```

### Good commits:

```
commit def456
"Fix: Correct momentum calculation in TrackReconstructor"

commit ghi789
"Docs: Add usage examples for ParticleAnalysis"

commit jkl012
"Refactor: Extract ParticleFilter into a separate class"
```

**Why?** Each commit should represent a single logical change.

# Why Atomic Commits Matter

---

## Enable `git bisect`

**Scenario:** A bug was introduced sometime in the last 100 commits

```
$ git bisect start
$ git bisect bad HEAD          # Current version is broken
$ git bisect good v1.2.3        # v1.2.3 was working
Bisecting: 50 revisions left to test

# Git automatically checks out a commit
$ ./run_tests.sh
$ git bisect good  # or bad

# Repeat ~7 times ...
# Git identifies the exact commit that broke it
```

This only works if commits are atomic and buildable!

# Commit Messages: Good vs Bad

✗ Bad:

```
"fix"  
"update"  
"stuff"  
"final version 2"  
"I hate git"
```

✓ Good:

Fix: Correct pT threshold  
  
Updated from 24 to 25 GeV  
to match the Run 3 trigger.  
  
JIRA: ATL-PHYS-2026-123

Why? Six months later, which one helps you understand what changed?

# Commit Message Rules

---

1. Separate subject from body with a blank line
2. Limit subject line to 50 characters
3. Capitalise the subject line
4. Do not end subject with a period
5. Use imperative mood in subject ("Fix bug" not "Fixed bug")
6. Wrap body at 72 characters
7. Use body to explain **what** and **why**, not **how**

## Template:

Short (50 chars or less) summary

More detailed explanatory text, if necessary. Wrap it to about 72 characters. Explain the problem this commit solves.

Fixes: #42

# Continuous Integration as a Factory

---

## Not Just Testing—Production

Modern CI does:

1.  **Build** the code (multiple compilers, platforms)
2.  **Test** functionality (unit, integration, system)
3.  **Analyse** code quality (static analysis, coverage)
4.  **Benchmark** performance (track regressions)
5.  **Publish** artifacts (plots, documentation, binaries)
6.  **Deploy** (if tests pass)

**Philosophy:** If CI passes, it's production-ready.

# Dual-Stack CI: GitHub + GitLab

---

## Why Both?

### GitHub Actions:

- Great for open-source projects
- Easy integration with GitHub features
- Large ecosystem of actions and industry standard

### GitLab CI (CERN Standard):

- Integrated with CERN infrastructure
- Access to CVMFS, EOS, LXPLUS runners
- Better for internal & self-hosted projects

**Best practice:** Master GitLab CI (what CERN uses), understand GitHub Actions (almost everywhere else)

# GitHub Actions: Basic CI

```
name: CI
on: [push, pull_request]

jobs:
  build:
    runs-on: ubuntu-latest
    container: rootproject/root:latest
    steps:
      - uses: actions/checkout@v4
      - run: cmake -B build -DCMAKE_BUILD_TYPE=Release
      - run: cmake --build build -j
      - run: ctest --test-dir build --output-on-failure
```

Add sanitisers as a separate job for cleaner output.

# GitLab CI (.gitlab-ci.yml)

```
stages: [build, test]

build:
  stage: build
  image: gitlab-registry.cern.ch/linuxsupport/alma9-base
  script:
    - source /cvmfs/sft.cern.ch/.../setup.sh
    - cmake -B build && cmake --build build
  artifacts:
    paths: [build/]

test:
  stage: test
  script:
    - ctest --test-dir build --output-on-failure
```

**Key:** Artifacts pass built files between stages.

# CI in Action

---

## The Workflow

1. **Student opens PR:** "Add electron selection criteria"
2. **CI automatically:**
  - Builds code
  - Runs tests
  - Compiles ROOT analysis macro
  - Generates histogram
  - Posts comment with plot link
3. **Reviewer sees:**
  - All checks passed
  - Visual comparison of electron distributions
  - Code coverage report
4. **Decision:** Merge with confidence

# Publishing Artifacts from CI

```
benchmark:  
  stage: test  
  script:  
    - ./run_benchmarks.sh > benchmark_results.txt  
    - python scripts/generate_plots.py  
  artifacts:  
    paths:  
      - benchmark_results.txt  
      - plots/  
    reports:  
      metrics: metrics.txt  
    expire_in: 1 week  
  only:  
    - main  
    - merge_requests
```

**Result:** Benchmark results and plots are available directly in the merge request for review.

# Code Review Best Practices

---

## For Authors

- Keep PRs small (<400 lines if possible)
- Write descriptive PR descriptions
- Link to relevant issues
- Respond to feedback promptly
- Don't take criticism personally

## For Reviewers

- Review within 24 hours
- Focus on correctness first, style second
- Ask questions, don't demand changes
- Approve when good enough, not perfect
- Use inline comments for specific issues

# Semantic Versioning

---

## MAJOR.MINOR.PATCH

**Example:** v2.3.1

- **MAJOR** (2): Incompatible API changes
- **MINOR** (3): Backwards-compatible new features
- **PATCH** (1): Backwards-compatible bug fixes

**Rules:**

- Start at v0.1.0 for initial development
- v1.0.0 means "production-ready with stable API"
- Increment MAJOR when breaking changes
- Increment MINOR when adding features
- Increment PATCH for bug fixes

**Learn more:** <https://semver.org/>

# CHANGELOG.md

---

## Keep Users Informed

### # Changelog

#### ## [2.3.1] - 2026-01-15

##### ### Fixed

- Corrected momentum calculation in TrackReconstructor (#142)
- Fixed memory leak in event loop (#145)

#### ## [2.3.0] - 2026-01-08

##### ### Added

- Support for Run 3 data format
- New electron identification algorithm

##### ### Changed

- Improved performance of track fitting by 30%

##### ### Deprecated

- Old track fitting API (will be removed in v3.0.0)

# Release Workflow

---

## Tagging and Distribution

```
# 1. Update version in code  
sed -i 's/VERSION "2.3.0"/VERSION "2.3.1"/' CMakeLists.txt  
# 2. Update CHANGELOG.md  
vim CHANGELOG.md  
  
# 3. Commit changes  
git add CMakeLists.txt CHANGELOG.md  
git commit -m "Release v2.3.1"  
# 4. Create an annotated tag  
git tag -a v2.3.1 -m "Version 2.3.1 - Bug fixes"  
  
# 5. Push tag  
git push origin v2.3.1  
  
# CI automatically builds and publishes release artifacts
```

# Summary: T&T L2

---

## Key Takeaways:

1. **Branching models** – Feature branches keep main stable
2. **Atomic commits** – Enable git bisect and clean history
3. **CI is production** – Automate build, test, analyse, publish
4. **Dual-stack CI** – GitHub Actions + GitLab CI
5. **Semantic versioning** – Communicate changes clearly
6. **CHANGELOG.md** – Document what changed and why

**Next:** Performance fundamentals – how to make code fast

# TOOLS & TECHNIQUES

---

## Lecture 3: Performance Fundamentals

"Premature optimisation is the root of all evil, but procrastination is worse"

60 minutes

Learn why code is slow and how to make it fast – scientifically

# T&T L3: Performance Fundamentals

---

## Why Performance Matters in HEP

Scale of the problem:

- LHC collisions: 40 MHz proton bunch crossings
- ATLAS/CMS: ~1 petabyte of data per second (before trigger)
- After trigger: ~1 GB/s to disk
- Full Run 3: ~100 petabytes per year

**Reality:** We can't store or process everything. Performance determines what science is possible.

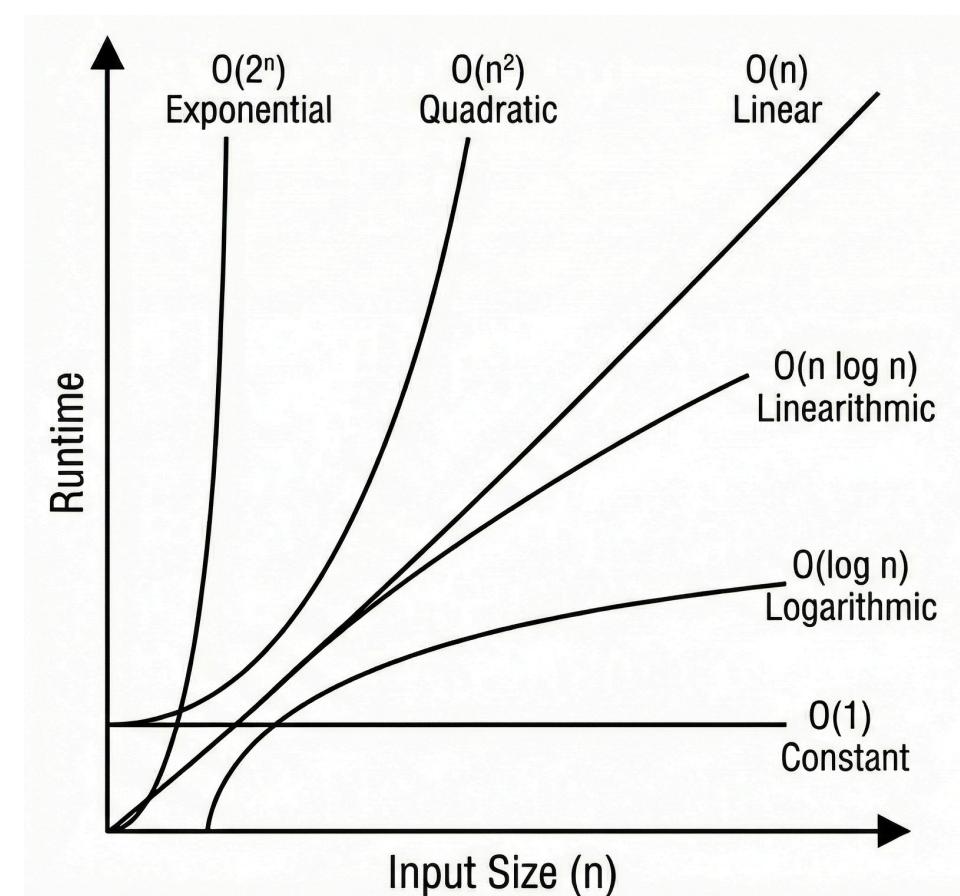
# Algorithmic Complexity

## Big-O Notation

**Definition:** How runtime scales with input size  $n$

Complexity	Name	Example
$O(1)$	Constant	Array access
$O(\log n)$	Logarithmic	Binary search
$O(n)$	Linear	Single loop
$O(n \log n)$	Linearithmic	Merge sort
$O(n^2)$	Quadratic	Nested loops
$O(2^n)$	Exponential	Naïve recursion

**Rule:** Always know the complexity of your algorithms



# HEP Example: Particle ID Lookup (1/5)

---

## ATLAS Use Case: Finding $Z \rightarrow ee$ Candidates

**Problem:** You have 10,000 electron candidates.

You need to find pairs with invariant mass near 91 GeV (Z boson).

**Bad: Linear Scan –  $O(n^2)$**

# HEP Example: Particle ID Lookup (2/5)

```
// Naïve approach: Check every pair
std::vector<ZCandidate> findZCandidates(const std::vector<Electron>& electrons) {
    std::vector<ZCandidate> candidates;

    for (size_t i = 0; i < electrons.size(); ++i) {
        for (size_t j = i+1; j < electrons.size(); ++j) {
            float mass = invariantMass(electrons[i], electrons[j]);
            if (std::abs(mass - 91.2) < 10.0) {
                candidates.push_back({electrons[i], electrons[j]});
            }
        }
    }
    return candidates;
}
// For 10,000 electrons: 50 million comparisons! Takes ~5 seconds
```

**Reality check:** ATLAS has millions of events. This is unusable.

# HEP Example: Particle ID Lookup (3/5)

Good: Sorted + Binary Search –  $O(n \log n)$

```
// Smart approach: Sort by pT, then search
std::vector<ZCandidate> findZCandidates(std::vector<Electron> electrons) {
    // Sort once:  $O(n \log n)$ 
    std::sort(electrons.begin(), electrons.end(),
              [](const Electron& a, const Electron& b) {
                  return a.pt() > b.pt();
              });

    std::vector<ZCandidate> candidates;
    // (continued on next slide)
```

# HEP Example: Particle ID Lookup (4/5)

Good: Sorted + Binary Search –  $O(n \log n)$

```
// (continued from previous slide)
// For each electron, binary search for compatible partners
for (size_t i = 0; i < electrons.size(); ++i) {
    auto range = findCompatibleRange(electrons, i);
    for (auto j = range.first; j != range.second; ++j) {
        // ... check invariant mass
    }
}
return candidates;
}
// For 10,000 electrons: ~100,000 comparisons. Takes ~0.1 seconds
// 50x faster!
```

# HEP Example: Particle ID Lookup (5/5)

---

## Best: Hash Map – O(1)

**Note:** When the problem allows keyed lookup (e.g., finding a particle by PDG ID), hash maps are ideal:

```
// Build index once
std::unordered_map<int, size_t> pdgIndex;
for (size_t i = 0; i < particles.size(); ++i) {
    pdgIndex[particles[i].pdgID()] = i;
}

// Lookup is constant time
int findParticleID(int pdgID) {
    auto it = pdgIndex.find(pdgID);
    return (it != pdgIndex.end()) ? it->second : -1;
}
// For 10,000 particles: ~1 comparison (amortised)
```

# The Memory Wall

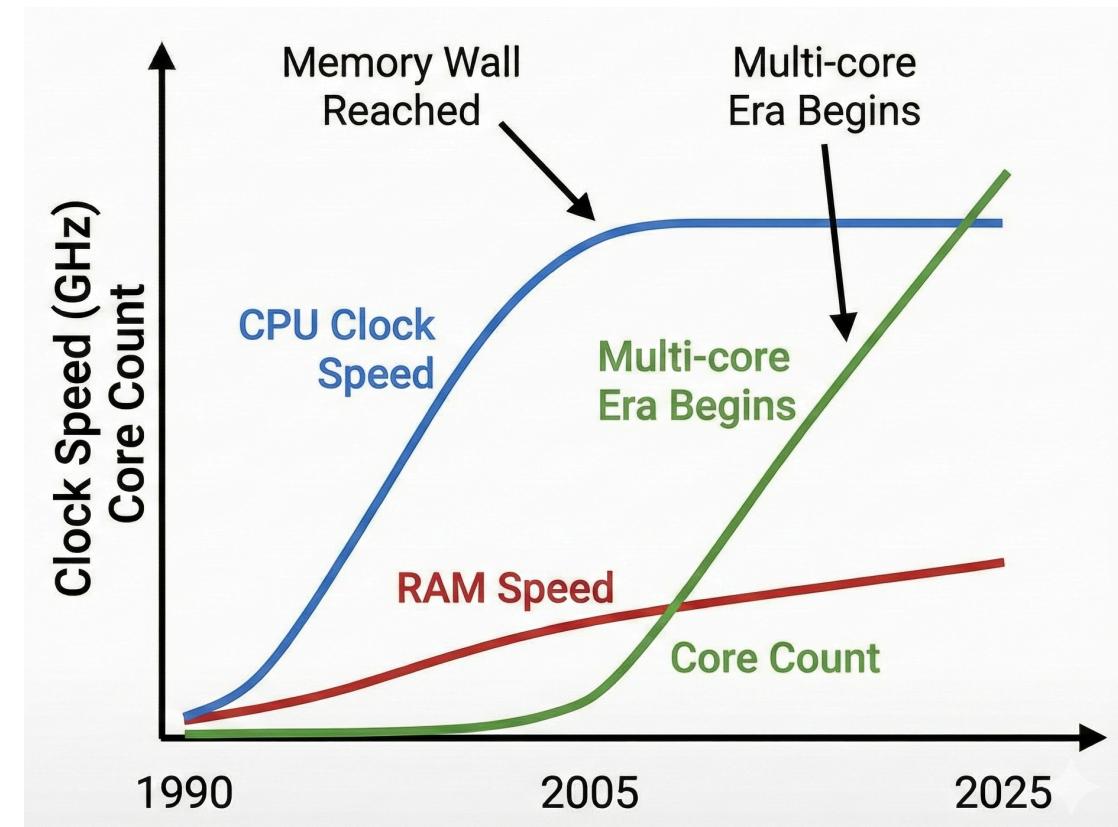
## Hardware Timeline

### Golden Age (1980–2005):

- 1990: 486 @ 50 MHz
- 2005: Pentium 4 @ 3.8 GHz
- CPU: 76× faster | RAM: 4× faster

### Modern Era (2005–now):

- 2025: Server CPU @ 3.5 GHz (up to 128 cores)
- Clock speed: barely changed
- Solution: More cores, not faster cores



# CPUs Are Fast, Memory Is Slow

---

Modern CPU (2026):

- L1 cache: ~4 cycles, 32–64 KB
- L2 cache: ~10–12 cycles, 256 KB–1 MB
- L3 cache: ~40–50 cycles, 8–64 MB
- Main RAM: ~200+ cycles

Gap: large!

Your algorithm might spend most of its time waiting for data.

ATLAS track reconstruction is often memory-bound in practice.

# Cache Lines and Locality (1/2)

---

## How Caches Work

**Cache line:** typically 64 bytes on x86-64. CPU fetches memory in chunks.

When you access `array[0]`, CPU loads `array[0 .. 15]` (if int = 4 bytes)

**Sequential access (fast):**

```
for (int i = 0; i < n; ++i) {
    sum += array[i]; // ✓ Cache hits
}
```

# Cache Lines and Locality (2/2)

---

## How Caches Work

Random access (slow):

```
for (int i = 0; i < n; ++i) {
    sum += array[randomIndex[i]]; // ✗ Cache misses
}
```

Sequential access can be an order of magnitude faster than random – often 10-20× on modern CPUs!

# Structure of Arrays vs Array of Structures

## The Critical Data Layout Choice

AoS (Array of Structures):

```
struct Particle {  
    float px, py, pz;  
    float energy;  
    int charge;  
};  
  
std::vector<Particle> particles;
```

Memory layout:

```
[px py pz E c][px py pz E c] ...
```

SoA (Structure of Arrays):

```
struct Particles {  
    std::vector<float> px;  
    std::vector<float> py;  
    std::vector<float> pz;  
    std::vector<float> energy;  
    std::vector<int> charge;  
};
```

Memory layout:

```
[px px px ... ][py py py ... ] ...
```

# Why SoA Matters

---

## Example benchmark (memory-bound loop)

**Task:** Calculate  $pT = \sqrt{(px^2 + py^2)}$  for 1M particles

```
// AoS: loads 20 bytes, uses 8
for (auto& p : particles) {
    pT = sqrt(p.px*p.px + p.py*p.py);
}
// Time: 45 ms

// SoA: loads only what's needed + vectorises
for (size_t i = 0; i < n; ++i) {
    pT = sqrt(px[i]*px[i] + py[i]*py[i]);
}
// Time: 6 ms
```

**Result:** ~7.5x speedup just from data layout in this scenario.

# Measuring Performance Safely

---

## Pitfalls to Avoid

Common mistakes:

1. **Frequency scaling:** CPU clock changes dynamically (Turbo Boost)
2. **Cold caches:** First run is always slower
3. **Background processes:** Other programs interfere
4. **Optimiser tricks:** Compiler might optimise away your benchmark
5. **Small inputs:** Doesn't represent real workload

**Solution:** Use proper benchmarking tools

# Google Benchmark

---

```
#include <benchmark/benchmark.h>

static void BM_PtCalc_AoS(benchmark::State& state) {
    auto particles = generateData(10000);
    for (auto _ : state) {
        float sum = 0;
        for (const auto& p : particles)
            sum += std::sqrt(p.px*p.px + p.py*p.py);
        benchmark::DoNotOptimize(sum);
    }
}
BENCHMARK(BM_PtCalc_AoS);
BENCHMARK_MAIN();
```

Compile: g++ -O3 -lbenchmark bench.cpp

# Using `perf` for System-Level Profiling

## Hardware Performance Counters

```
# Measure cache misses
$ perf stat -e cache-references,cache-misses,cycles,instructions \
    ./reconstruction input.root

Performance counter stats:
      234,567,890  cache-references
      45,678,901  cache-misses          # 19.5% of all cache refs
      1,234,567,890  cycles
      2,345,678,901  instructions        # 1.90  insns per cycle

      2.345678 seconds time elapsed
```

- High cache miss rate (>10%) → memory-bound
- Low instructions per cycle (<1.5) → CPU stalls

# Modern HEP: ROOT RDataFrame

---

## Writing Thread-Safe Analysis Code

Traditional ROOT (single-threaded):

```
TFile* file = TFile::Open("data.root");
TTree* tree = (TTree*)file->Get("Events");

float pt;
tree->SetBranchAddress("pt", &pt);
TH1F* h = new TH1F("h", "pT distribution", 100, 0, 200);

for (Long64_t i = 0; i < tree->GetEntries(); ++i) {
    tree->GetEntry(i);
    if (pt > 20.0) h->Fill(pt);
}
```

**Problem:** Sequential processing – can't use multiple cores

# ROOT RDataFrame + IMT

---

```
#include <ROOT/RDataFrame.hxx>

// Enable multithreading BEFORE creating RDataFrame
ROOT::EnableImplicitMT();

ROOT::RDataFrame df("Events", "data.root");
auto h = df.Filter("pt > 20.0")
           .Histo1D({"h", "pT", 100, 0, 200}, "pt");
h→Draw();
```

**Result:** Often 2–8× speedup on modern CPUs (I/O can limit).

# Demo: RDataFrame Speedup

---

```
// Single-threaded  
ROOT::DisableImplicitMT();  
benchmark(); // Time: 8.2s  
  
// Multi-threaded (8 cores)  
ROOT::EnableImplicitMT(8);  
benchmark(); // Time: 1.4s → 5.9× speedup!
```

Typical results:

Threads	Time	Speedup
1	8.2s	1.0×
4	2.3s	3.6×
8	1.4s	5.9×

**One line of code** – no manual thread management needed!

# Summary: T&T L3

---

## Key Takeaways:

1. **Algorithm choice matters most** –  $O(n)$  vs  $O(\log n)$  vs  $O(1)$
2. **Memory is the bottleneck** – Cache-friendly code is fast code
3. **SoA > AoS for HEP** – Better cache usage and vectorisation
4. **Measure correctly** – Use Google Benchmark and perf
5. **ROOT RDataFrame** – Easy thread safety with implicit MT

**Next:** Hands-on exercises to apply these concepts

# TOOLS & TECHNIQUES

---

## Exercises: E1 + E2

"Learning by doing — the only way it sticks"

2 hours total

Pairs, hands-on, real HEP data

Debug broken code with sanitisers  
Optimise real ATLAS Open Data I/O

# Exercise E1: Make It Right

---

## Debugging with Sanitisers (60 minutes)

**Scenario:** You've inherited a "broken" track reconstruction library. It compiles but produces incorrect results and occasionally crashes.

**Your mission:**

1. Use **ASan** to find memory bugs
2. Write **Catch2 tests** to prevent regression
3. Fix the bugs
4. Ensure all tests pass in CI

**Files provided:** `reconstruction.cpp`, `reconstruction.h`, `test_reconstruction.cpp`,  
`CMakeLists.txt`

# E1: Getting Started

---

## Setup

```
# Clone the exercise repository
git clone https://github.com/artfisica/csc2026.git
cd exercise-e1

# Create development container (optional but recommended)
devcontainer open .

# Build with ASan
cmake -B build -S . \
    -DCMAKE_BUILD_TYPE=Debug \
    -DCMAKE_CXX_FLAGS="-fsanitize=address -g"
cmake --build build

# Run tests
cd build && ctest --output-on-failure
```

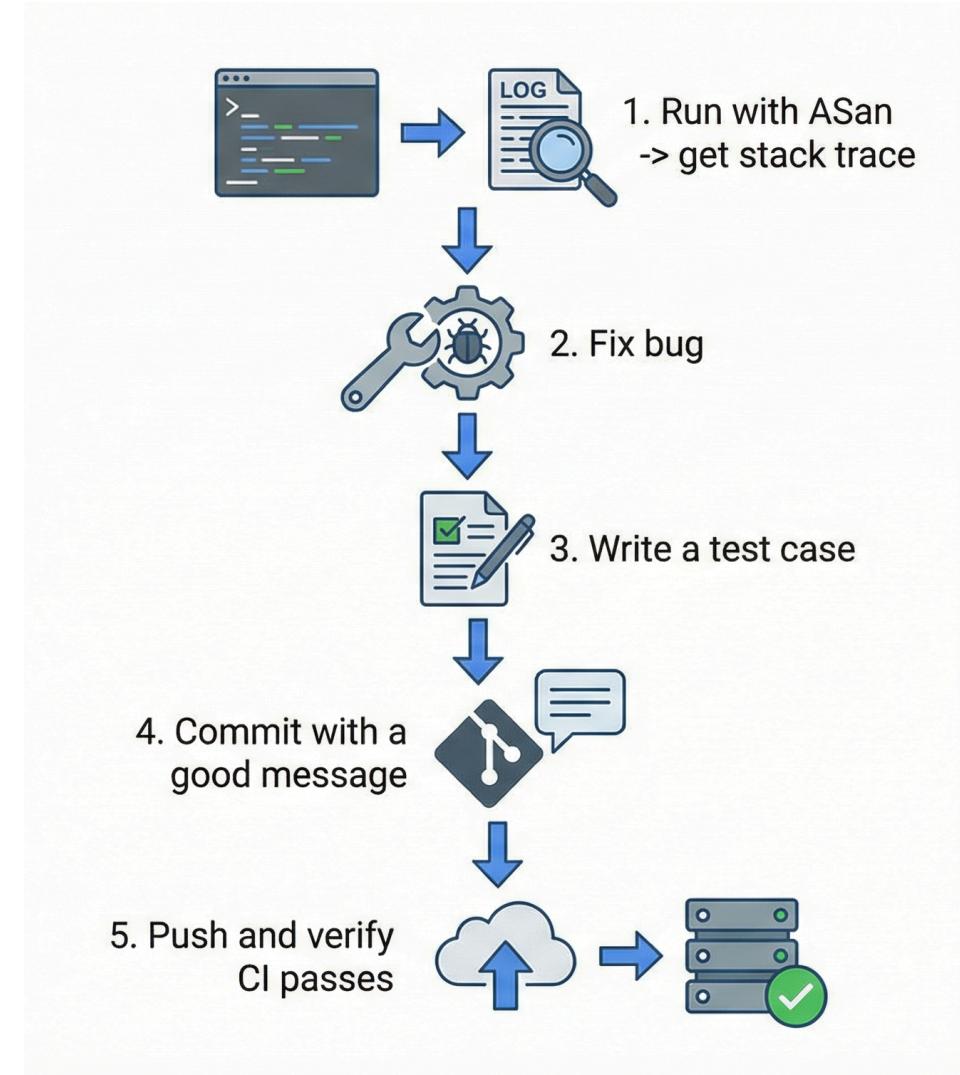
# E1: Expected Bugs

You'll encounter:

1. **Buffer overflow:** Array access out of bounds
2. **Use-after-free:** Accessing deleted memory
3. **Memory leak:** Forgetting to deallocate

Workflow:

1. Run with ASan → get stack trace
2. Fix bug
3. Write a test case that would catch it
4. Commit with a good message
5. Push and verify CI passes



# Exercise E2: Make It Fast Enough (1/2)

---

## Performance Optimisation (60 minutes)

### Two tracks available

Track A (Algorithmic): Optimise a particle sorting/search routine

- Profile with `perf`
- Identify hotspots
- Apply algorithmic improvements
- Demonstrate speedup

## Exercise E2: Make It Fast Enough (2/2)

---

Track B (Real I/O with ATLAS Open Data): Optimise ROOT file reading

- Profile I/O throughput
- Optimise branch reading
- Test RDataFrame
- Generate CI artifact showing MB/s improvement

# E2 Track A: Algorithmic Optimisation

---

## The Challenge

**File:** `particle_search.cpp`

**Current implementation:** Naïve  $O(n^2)$  algorithm for finding particle pairs

**Your tasks:**

1. Profile with Google Benchmark
2. Identify the complexity bottleneck
3. Implement  $O(n \log n)$  solution
4. Verify correctness with tests
5. Document speedup in CI artifact

**Hint:** Think about sorting + binary search, or spatial indexing

# E2 Track B: ROOT I/O Optimisation

---

## The Challenge

**Files:** Real ATLAS Open Data (13 TeV) – suitable for teaching derived from Run formats ([the Higgs discovery in 2012 used 7-8 TeV data](#))

**Dataset:** NanoAOD-style ROOT files with electron, muon, jet collections (~100 MB each)

**Current code:** Reads all branches, slow iteration

**Your mission:**

1. **Baseline:** Measure current I/O throughput (MB/s)
2. **Optimise:** Use `SetBranchStatus` to disable unneeded branches
3. **Optimise:** Adjust basket size/cache settings
4. **Modern approach:** Try RDataFrame
5. **Deliverable:** CI artifact with before/after throughput comparison

# E2 Track B: Getting Started

```
// Baseline measurement
void baseline() {
    TFile* file = TFile::Open("data.root");
    TTree* tree = (TTree*)file->Get("Events");
    // Read all branches (slow)
    tree->SetBranchStatus("*", 1);

    Long64_t nEntries = tree->GetEntries();
    auto start = std::chrono::high_resolution_clock::now();

    for (Long64_t i = 0; i < nEntries; ++i) {
        tree->GetEntry(i);
        // Process ...
    }

    auto end = std::chrono::high_resolution_clock::now();
    // Calculate MB/s ...
}
```

# E2 Track B: Optimisation Strategies (1/2)

## 1. Selective Branch Reading

```
// Only read branches you need
tree→SetBranchStatus("*", 0);           // Disable all
tree→SetBranchStatus("Electron_pt", 1); // Enable specific
tree→SetBranchStatus("Electron_eta", 1);
tree→SetBranchStatus("Electron_phi", 1);
```

## 2. Cache Tuning

```
// Increase cache size (default is often too small)
tree→SetCacheSize(50 * 1024 * 1024); // 50 MB
tree→AddBranchToCache("Electron_*", kTRUE);
```

## E2 Track B: Optimisation Strategies (2/2)

---

### 3. RDataFrame (Best Practice)

```
ROOT::RDataFrame df("Events", "data.root");
auto filtered = df.Filter("Electron_pt[0] > 25");
```

# Both Tracks: CI Integration

---

## Your CI Should:

1. **Build** with optimisations enabled (-03)
2. **Run benchmarks** automatically
3. **Generate comparison** (before/after)
4. **Post artifact** (plot or table)

Example `.gitlab-ci.yml` addition:

```
benchmark:  
  stage: test  
  script:  
    - ./build/benchmark_baseline > baseline.txt  
    - ./build/benchmark_optimised > optimised.txt  
    - python scripts/compare.py baseline.txt optimised.txt  
  artifacts:  
    paths:  
      - comparison.png  
      - improvement_report.txt
```

# Exercise Tips & Getting Help

---

## Working Together

- Work in pairs if you can
- Talk through your approach
- Commit often with good messages
- Test as you go
- Ask for help, seriously
- Document what you tried

**When stuck (totally normal!):**

1. Talk to partner first
2. Read error messages
3. Raise your hand
4. Ask the group
5. Use the tools (ASan/TSan/perf)

**There are no stupid questions!** Keep notes of what worked.

# Time Management

---

## E1: "Make it Right" (~60 min)

- 20 min: Find bugs with ASan
- 20 min: Write tests
- 15 min: Fix bugs
- 5 min: CI verification

## E2: "Make it Fast" (~60 min)

- 15 min: Baseline
- 20 min: Find bottleneck
- 20 min: Optimise
- 5 min: Verify

## Tips

- Get it working, then improve.
- Prioritise correctness first

# SOFTWARE DESIGN

---

## Lecture 1: Many-Core Foundations

"The free lunch is over — welcome to the parallel era"

60 minutes

Understand why and how to write parallel code

Power wall, Amdahl's Law, OpenMP, TBB

# SD L1: Many-Core Foundations

---

## The Hardware Reality

### 2005: The Industry's Biggest Pivot

Before: "Wait 18 months, get 2x faster CPU"

After: "Here are more cores. Figure it out."

#### What happened:

- Intel/AMD cancel 10 GHz plans
- Power  $\approx C \times V^2 \times f$  (raising  $f$  often requires higher  $V$  → superlinear power growth)
- Result: Clock speeds stuck at ~3-5 GHz

Everything about "faster computers = faster code" stopped being true.

# Today's Landscape

---

**Consumer:** 8-32 cores typical

**HEP (WLCG):** 32-128 cores per node

**Total WLCG:** ~1 million cores worldwide

- <https://home.cern/science/computing/grid>

**The challenge:** Your code must use all cores or waste 90%+ compute!

# The Power Wall

---

## Why Clock Speeds Stopped Growing

Power consumption:  $P \approx C \times V^2 \times f$

Where:

- $C$  = capacitance
- $V$  = voltage
- $f$  = frequency (clock speed)

**Problem:** Power grows with frequency. Higher clocks → more heat → can't cool

**Solution:** More cores at lower frequency instead of fewer cores at higher frequency

**Implication:** To use modern hardware, **we must parallelise**

# Amdahl's Law

---

## The Serial Bottleneck

**Formula:** Speedup =  $1 / [(1 - P) + P/N]$  | Where P = parallel fraction, N = cores

Cores	Speedup	Time	Efficiency
1	1.0x	100s	100%
2	1.75x	57s	88%
4	2.96x	34s	74%
8	4.21x	24s	53%
$\infty$	6.67x	15s	0%

**ATLAS Example:** 100s/event total:

- I/O: 15s (serial);
- Track/vertex: 85s (parallel);
- P = 0.85

**Lesson:** 15% serial code limits max speedup to 6.67x!

# Gustafson's Law

---

## Scale the Problem, Not the Solution

**Alternative view:** As hardware improves, we tackle larger problems

**Formula:** Speedup =  $N - (1 - P) \times (N - 1)$

Where P is a parallel fraction, N is the number of cores

**Key insight:** If you can scale problem size, parallel efficiency improves

**HEP context:** We don't process the same 1M events faster — we process 1B events

# Concurrency vs Parallelism

---

## Important Distinction

### Concurrency

- About **structure**
- Multiple tasks in progress
- Doesn't require multiple cores
- Example: Event-driven I/O

### Parallelism

- About **execution**
- Multiple tasks running simultaneously
- Requires multiple cores
- Example: Multi-threaded processing

In HEP: We need both

- Concurrency: Handle multiple data streams
- Parallelism: Process events on multiple cores

# Parallel Programming: Rules of Thumb

---

## ✓ Good candidates:

- Loop with >10k independent items
- Each iteration >1 microsecond
- Zero dependencies between iterations
- Contiguous arrays (not lists/trees)

## ✗ Bad candidates:

- <1k items (overhead > benefit)
- Heavy synchronisation needed
- Rarely executed code

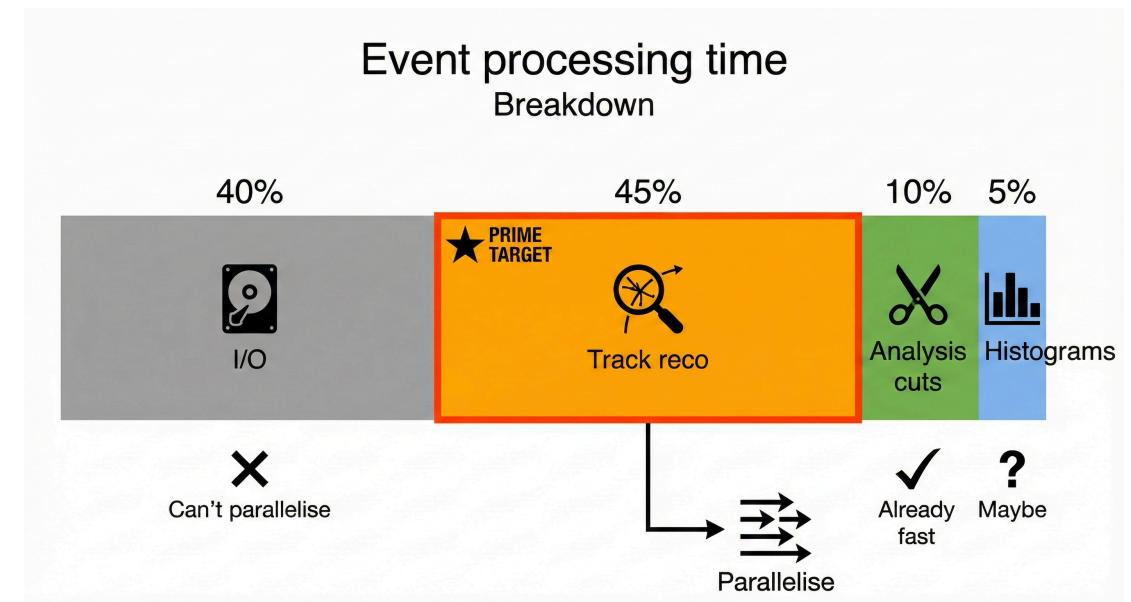
**80/20 rule:** Profile first! 80% of time in 20% of code.

# Real LHC experiment Breakdown

## Event processing time:

- 40% I/O → Can't parallelise
- 45% Track reco → **PRIME TARGET**
- 10% Analysis cuts → Already fast
- 5% Histograms → Maybe

**Focus:** Parallelise the 45%



# C++ Standard Library Threading

```
#include <thread>
#include <vector>

void processEvent(int eventID) {
    // Heavy computation
}

int main() {
    std::vector<std::thread> workers;
    for (int i = 0; i < 8; ++i) {
        workers.emplace_back(processEvent, i);
    }
    for (auto& t : workers) {
        t.join(); // Wait for completion
    }
}
```

**Problem:** Manual thread management is error-prone

# Atomics: Lock-Free Synchronisation (1/2)

## Safe Counter Without Locks

```
#include <atomic>
#include <thread>
#include <vector>

std::atomic<int> eventCount{0};

void processEvents(int n) {
    for (int i = 0; i < n; ++i) {
        // Do work ...
        eventCount.fetch_add(1, std::memory_order_relaxed);
    }
}
```

# Atomics: Lock-Free Synchronisation (2/2)

---

## Safe Counter Without Locks

```
int main() {
    std::vector<std::thread> workers;
    for (int i = 0; i < 8; ++i) {
        workers.emplace_back(processEvents, 1000);
    }

    for (auto& t : workers) t.join();
    std::cout << "Total events: " << eventCount << std::endl;
}
```

# Mutexes: When You Must Synchronise (1/2)

## Protecting Shared State

```
#include <mutex>
#include <map>
#include <thread>

std::map<int, Result> resultsCache;
std::mutex cacheMutex;

void storeResult(int eventID, const Result& result) {
    std::lock_guard<std::mutex> lock(cacheMutex);
    resultsCache[eventID] = result; // Protected
}
// (continued on next slide)
```

# Mutexes: When You Must Synchronise (2/2)

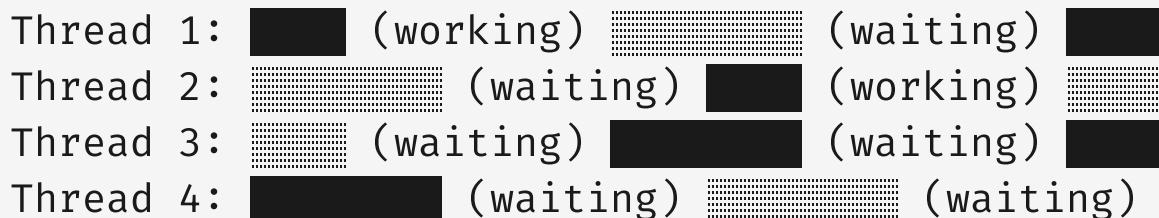
## Protecting Shared State

```
// (continued from previous slide)
Result loadResult(int eventID) {
    std::lock_guard<std::mutex> lock(cacheMutex);
    return resultsCache.at(eventID);
}
```

**Warning:** Mutexes are slow. Avoid when possible. Use lock-free algorithms or partition data instead.

# Why Avoid Mutexes? | Contention Kills Performance

**Problem:** Multiple threads waiting for the same lock



**Result:** Threads spend more time waiting than working

Better approaches:

1. **Partition data** – Give each thread independent data
2. **Lock-free structures** – Use atomics or lock-free queues
3. **Thread-local storage** – Each thread has own copy

# OpenMP: Easy Parallelism

```
#include <omp.h>

void parallelProcess(std::vector<Event>& events) {
    #pragma omp parallel for
    for (size_t i = 0; i < events.size(); ++i) {
        events[i].reconstruct();
    }
}
```

Compile: g++ -fopenmp program.cpp

Control threads: OMP\_NUM\_THREADS=8 ./program

# OpenMP: Reductions

## Parallel Aggregation

```
double sumEnergy(const std::vector<Particle>& particles) {
    double total = 0.0;

    #pragma omp parallel for reduction(+:total)
    for (size_t i = 0; i < particles.size(); ++i) {
        total += particles[i].energy();
    }

    return total;
}
```

**How it works:** Each thread maintains a local sum, then combines at the end

**Supported operations:** +, \*, -, &, |, ^, &&, ||, min, max

# oneTBB: Task-Based Parallelism

```
#include <tbb/parallel_for.h>
#include <tbb/blocked_range.h>

void process(std::vector<Event>& events) {
    tbb::parallel_for(
        tbb::blocked_range<size_t>(0, events.size()),
        [&](const auto& r) {
            for (size_t i = r.begin(); i != r.end(); ++i)
                events[i].reconstruct();
        }
    );
}
```

**Advantages:** Work stealing, composable, modern C++ API

# TBB: Parallel Reduction

---

```
#include <tbb/parallel_reduce.h>

double sumEnergy(const std::vector<Particle>& p) {
    return tbb::parallel_reduce(
        tbb::blocked_range<size_t>(0, p.size()),
        0.0,
        [&](const auto& r, double s) {
            for (size_t i = r.begin(); i != r.end(); ++i)
                s += p[i].energy();
            return s;
        },
        std::plus<double>()
    );
}
```

# HEP Context: Where Threading Is Used (1/2)

---

## CMSSW (CMS Software)

- Framework: TBB-based task scheduling
- Event-level parallelism: Multiple events processed concurrently
- Module-level parallelism: Independent modules run in parallel

## Geant4 (Detector Simulation)

- Multi-threading mode: Multiple events simulated simultaneously
- Thread-local geometry: Each thread has navigation state
- Shared read-only data: Geometry, physics processes

# HEP Context: Where Threading Is Used (2/2)

---

## ROOT

- **RDataFrame**: Implicit multithreading
- **TTree::GetEntry**: Can be parallelised with care
- **Histogram filling**: Thread-safe with proper setup

# Common Threading Pitfalls (1/2)

## 1. Race Conditions

```
int counter = 0;  
#pragma omp parallel for  
for (int i = 0; i < 1000; ++i) {  
    counter++; // WRONG, Multiple threads writing  
}  
// Solution: Use reduction or atomic
```

## 2. False Sharing

```
struct Data { int count; }; // Only 4 bytes  
Data perThread[8]; // Each thread's data  
// Problem: All in the same cache line → cache ping-pong!  
// Solution: struct Data { int count; char padding[60]; };
```

# Common Threading Pitfalls (2/2)

---

## 3. Deadlocks

```
// Thread 1: lock(A), then lock(B)  
// Thread 2: lock(B), then lock(A) → DEADLOCK!  
// Fix: Always acquire locks in the same order
```

# Parallel Debugging Checklist

---

When parallel code fails:

1.  Does it work with 1 thread? (If no → not a parallel bug)
2.  Run with TSan
3.  Check false sharing (variables <64 bytes apart?)
4.  Verify reduction correctness
5.  Check load balance
6.  Inspect scaling ( $2\times$  cores  $\approx 2\times$  speedup?)

**90% of bugs caught by TSan + this list!**

# Summary: SD L1

---

## Key Takeaways:

1. **Power wall** – Clock speeds stopped; we must go wide
2. **Amdahl's Law** – Serial sections limit speedup
3. **Primitives** – Threads, atomics, mutexes (avoid)
4. **OpenMP** – Easy pragma-based parallelism
5. **TBB** – Task-based, composable, better load balancing
6. **HEP frameworks** – CMSSW, Geant4, ROOT all use threading

**Next:** Data layout and parallel patterns

# SOFTWARE DESIGN

---

## Lecture 2: Patterns & Data Layout

"How you organise your data matters more than your algorithm"

60 minutes

Master parallel patterns and cache-friendly programming

Map/Reduce, SoA vs AoS, vectorisation, false sharing

# SD L2: Patterns & Data Layout

---

## Why Patterns Matter

Common parallel problems have common solutions

Instead of reinventing wheels:

1. Learn proven patterns
2. Recognise which pattern fits your problem
3. Apply the pattern correctly
4. Avoid common pitfalls

Today's patterns:

- Map/Reduce
- Pipeline
- Tiling

# Map/Reduce Pattern (1/2)

---

## The Most Common Parallel Pattern

**Map:** Apply function to each element independently

```
// Sequential
for (auto& x : data) {
    x = transform(x);
}

// Parallel (map)
#pragma omp parallel for
for (size_t i = 0; i < data.size(); ++i) {
    data[i] = transform(data[i]);
}
```

# Map/Reduce Pattern (2/2)

---

## The Most Common Parallel Pattern

**Reduce:** Combine results into a single value

```
// Parallel (reduce)
double sum = 0.0;
#pragma omp parallel for reduction(+:sum)
for (size_t i = 0; i < data.size(); ++i) {
    sum += data[i];
}
```

# Map/Reduce in HEP (1/2)

## Event Processing

```
// Map: Process each event independently
std::vector<Result> results(events.size());

tbb::parallel_for(
    tbb::blocked_range<size_t>(0, events.size()),
    [&](const tbb::blocked_range<size_t>& r) {
        for (size_t i = r.begin(); i != r.end(); ++i) {
            results[i] = processEvent(events[i]);
        }
    }
);
```

# Map/Reduce in HEP (2/2)

## Event Processing

```
// Reduce: Count events passing selection
int passing = tbb::parallel_reduce(
    tbb::blocked_range<size_t>(0, results.size()),
    0,
    [&](const tbb::blocked_range<size_t>& r, int count) {
        for (size_t i = r.begin(); i != r.end(); ++i) {
            if (results[i].passes()) count++;
        }
        return count;
},
std::plus<int>()
);
```

# Pipeline Pattern

---

## Stages of Processing

**Idea:** Divide work into sequential stages, parallelise within stages

```
Input → [Stage 1] → [Stage 2] → [Stage 3] → Output  
      ↓ ↓ ↓      ↓ ↓ ↓      ↓ ↓ ↓  
      parallel    parallel    parallel
```

**HEP Example:**

1. **Stage 1:** Read events from file
2. **Stage 2:** Reconstruct tracks/clusters (parallel)
3. **Stage 3:** Apply physics selection (parallel)
4. **Stage 4:** Fill histograms (need synchronisation)

# Pipeline with TBB (1/2)

---

```
#include <tbb/pipeline.h>

void processPipeline(const std::string& inputFile) {
    tbb::parallel_pipeline(
        /*max_tokens=*/16,
        tbb::make_filter<void, Event*>(
            tbb::filter_mode::serial_in_order,
            [&](tbb::flow_control& fc) → Event* {
                Event* evt = readNextEvent(inputFile);
                if (!evt) fc.stop();
                return evt;
            }
        ) &
    // (continued on next slide)
```

# Pipeline with TBB (2/2)

```
// (continued from previous slide)
tbb::make_filter<Event*, Event*>(
    tbb::filter_mode::parallel,
    [](Event* evt) {
        evt->reconstruct();
        return evt;
    }
) &
tbb::make_filter<Event*, void>(
    tbb::filter_mode::parallel,
    [](Event* evt) {
        evt->select();
        delete evt;
    }
);
}
```

# Tiling Pattern

---

## Block-Based Processing

**Idea:** Divide the problem into tiles/blocks that fit in the cache

T1	T2	T3
T4	T5	T6
T7	T8	T9

Process one tile at a time  
→ Better cache reuse  
→ Improved locality

## Applications:

- Matrix operations & Image processing
- Grid-based simulations

# Tiling Example: Matrix Multiplication (1/2)

---

```
// Naïve (cache-unfriendly)
for (int i = 0; i < N; ++i) {
    for (int j = 0; j < N; ++j) {
        for (int k = 0; k < N; ++k) {
            C[i][j] += A[i][k] * B[k][j];
        }
    }
}
```

# Tiling Example: Matrix Multiplication (2/2)

```
// Tiled (cache-friendly)
const int TILE = 64;
for (int ii = 0; ii < N; ii += TILE) {
    for (int jj = 0; jj < N; jj += TILE) {
        for (int kk = 0; kk < N; kk += TILE) {
            for (int i = ii; i < std::min(ii+TILE, N); ++i) {
                for (int j = jj; j < std::min(jj+TILE, N); ++j) {
                    for (int k = kk; k < std::min(kk+TILE, N); ++k) {
                        C[i][j] += A[i][k] * B[k][j];
                    }
                }
            }
        }
    }
}
```

# Data Layout: AoS vs SoA Revisited

---

## Why This Matters for Parallelism

Array of Structures (AoS):

```
struct Particle {  
    float px, py, pz, E;  
    int pdgID;  
};  
std::vector<Particle> particles;
```

Problems:

1. Cache line waste when accessing a single field
2. Cannot vectorise easily
3. False sharing between threads

# Data Layout: SoA Benefits

---

Structure of Arrays (SoA):

```
struct Particles {  
    std::vector<float> px, py, pz, E;  
    std::vector<int> pdgID;  
    size_t size() const { return px.size(); }  
};
```

Benefits:

1.  Contiguous memory for each field
2.  Auto-vectorisation by compiler
3.  No false sharing (different cache lines)
4.  Better cache utilisation

Trade-off: More complex indexing, but worth it

# False Sharing Example

---

## The Hidden Performance Killer

```
// Bad: False sharing
struct ThreadData {
    int counter; // Only 4 bytes
};
ThreadData data[8]; // Array of per-thread data

#pragma omp parallel
{
    int tid = omp_get_thread_num();
    for (int i = 0; i < 1000000; ++i) {
        data[tid].counter++; // Looks independent, but ...
    }
}
```

**Problem:** `counter`s share cache lines. Writing from one thread invalidates the cache for others.

# False Sharing: Solution

---

## Padding to Cache Line Boundaries

```
// Good: Pad to cache line size (64 bytes)
struct ThreadData {
    int counter;
    char padding[60]; // Total: 64 bytes
};
ThreadData data[8]; // Now each on different cache line

// Or use C++17:
struct alignas(64) ThreadData {
    int counter;
};
```

**Result:** Each thread's data on separate cache line → no false sharing

**Rule:** Per-thread data should be at least 64 bytes apart.

# Vectorisation (SIMD)

---

## Single Instruction, Multiple Data

**Modern CPUs:** Execute the same operation on 4–16 values simultaneously

AVX-512 Example:

```
// Scalar (one at a time)
for (int i = 0; i < n; ++i) {    c[i] = a[i] + b[i];    }
// Vectorised (16 floats at once)
for (int i = 0; i < n; i += 16) {
    __m512 va = _mm512_loadu_ps(&a[i]);
    __m512 vb = _mm512_loadu_ps(&b[i]);
    __m512 vc = _mm512_add_ps(va, vb);
    _mm512_storeu_ps(&c[i], vc);
}
```

**Speedup:** Up to 16x for simple operations

# Auto-Vectorisation (1/2)

---

## Let the Compiler Do It

Most compilers auto-vectorise simple loops:

```
// This loop can be auto-vectorised
void add(float* a, float* b, float* c, size_t n) {
    for (size_t i = 0; i < n; ++i) {
        c[i] = a[i] + b[i];
    }
}
```

# Auto-Vectorisation (2/2)

---

## Let the Compiler Do It

Compile with:

```
g++ -O3 -march=native -ftree-vectorize -fopt-info-vec
```

Requirements for auto-vectorisation:

- Contiguous arrays (not lists or trees)
- No dependencies between iterations
- Aligned data (64-byte alignment helps)

# SoA Enables Vectorisation

```
// AoS: Hard to vectorise
struct Particle { float px, py, pz, E; };
std::vector<Particle> particles;

for (size_t i = 0; i < particles.size(); ++i) {
    particles[i].px *= scale; // Compiler struggles here
}

// SoA: Easy to vectorise
struct Particles {
    std::vector<float> px, py, pz, E;
};

for (size_t i = 0; i < particles.px.size(); ++i) {
    particles.px[i] *= scale; // Compiler auto-vectorises
}
```

# HEP Framework Scheduling

---

## Task-Based Parallelism in CMSSW

### Traditional (Run 1-2):

- Sequential module execution
- One event at a time

### Modern (Run 3+):

- Task graph of module dependencies
- Multiple events in flight
- Modules run when inputs are ready

```
Event 1: [Read] → [Unpack] → [Reco] → [Select] → [Write]  
Event 2:           [Read] → [Unpack] → [Reco] → [Select]  
Event 3:           [Read] → [Unpack] → [Reco]
```

**Key:** Modules with no dependencies run in parallel

# Summary: SD L2

---

## Key Takeaways:

1. **Patterns** – Map/Reduce, Pipeline, Tiling
2. **Data layout** – SoA beats AoS for parallelism
3. **False sharing** – Pad per-thread data to 64 bytes
4. **Vectorisation** – SIMD gives 4–16× speedup
5. **HEP frameworks** – Task-based scheduling in CMSSW

**Next:** Profiling and optimisation workflow

# SOFTWARE DESIGN

---

## Lecture 3: Profiling to Optimisation

"In God we trust; all others must bring data"

60 minutes

Profile scientifically, optimise systematically

perf, flame graphs, roofline model, real case study

# SD L3: Profiling to Optimisation

---

## The Optimisation Workflow

### Four-Step Process:

1. **Measure** – Profile to find hotspots
2. **Analyse** – Understand why it's slow
3. **Optimise** – Apply targeted improvements
4. **Verify** – Confirm speedup and correctness

**Golden Rule:** Never optimise without profiling first

**Corollary:** Profile after every change to verify improvement

# Why Profile First?

---

## Avoid Wasted Effort

### Without profiling:

- Optimise the wrong code
- Spend days on 1% hotspot
- Make code complex for no gain
- Miss the real bottleneck

### With profiling:

- Focus on 80% of runtime
- Quick wins are visible immediately
- Quantify every improvement
- Understand trade-offs

**Example:** Developer spends a week optimising a loop that takes 0.5% of runtime. Meanwhile, I/O takes 60%.

# Profiling Tools

---

## The Essential Toolkit

Tool	Purpose	Platform
<code>perf</code>	System-level CPU profiling	Linux
<code>gprof</code>	Function-level profiling	Linux, macOS
<code>Valgrind/Callgrind</code>	Detailed call graphs	Linux
<code>Intel VTune</code>	Advanced CPU/memory analysis	Intel CPUs
<code>Google Benchmark</code>	Micro-benchmarking	Cross-platform
<code>ROOT TStopwatch</code>	HEP-specific timing	Cross-platform

**Today's focus:** `perf` (most common in HEP)

# Using perf: Basic Commands

---

## CPU Profiling

```
# Record profile (samples every ~100ms)
$ perf record -g ./reconstruction input.root

# View results interactively
$ perf report

# Or generate text report
$ perf report --stdio > profile.txt

# Annotate specific function
$ perf annotate processEvent

# Compare two runs
$ perf diff before.data after.data
```

### Key options:

- `-g`: Capture call graphs
- `-F 1000`: Sample at 1000 Hz
- `-e cycles`: Profile CPU cycles (default)

# Reading perf Output

## Example Report

#	Overhead	Command	Shared Object	Symbol
#	.....	.....	.....	.....
	45.23%	reco	libreco.so	[.] TrackFitter::fit()
	18.67%	reco	libreco.so	[.] ClusterBuilder::build()
	12.34%	reco	libRoot.so	[.] TTree ::GetEntry()
	8.91%	reco	libreco.so	[.] ParticleID::classify()
	5.43%	reco	libc.so.6	[.] memcpy
	3.21%	reco	libreco.so	[.] EnergyCorrection::apply()
	...			

- 45% of time in `TrackFitter::fit()` → Primary target
- 18% in `ClusterBuilder::build()` → Secondary target
- 12% in ROOT I/O → Consider optimising file reading

# Flame Graphs

---

## Visualising Call Stacks

Tool: [FlameGraph](#) by Brendan Gregg

```
# Generate flame graph
$ perf record -F 99 -g ./reconstruction input.root
$ perf script | stackcollapse-perf.pl | flamegraph.pl > flame.svg
$ firefox flame.svg
```

### How to read:

- Width = % of samples
- Stack depth = call hierarchy
- Click to zoom
- Search for function names

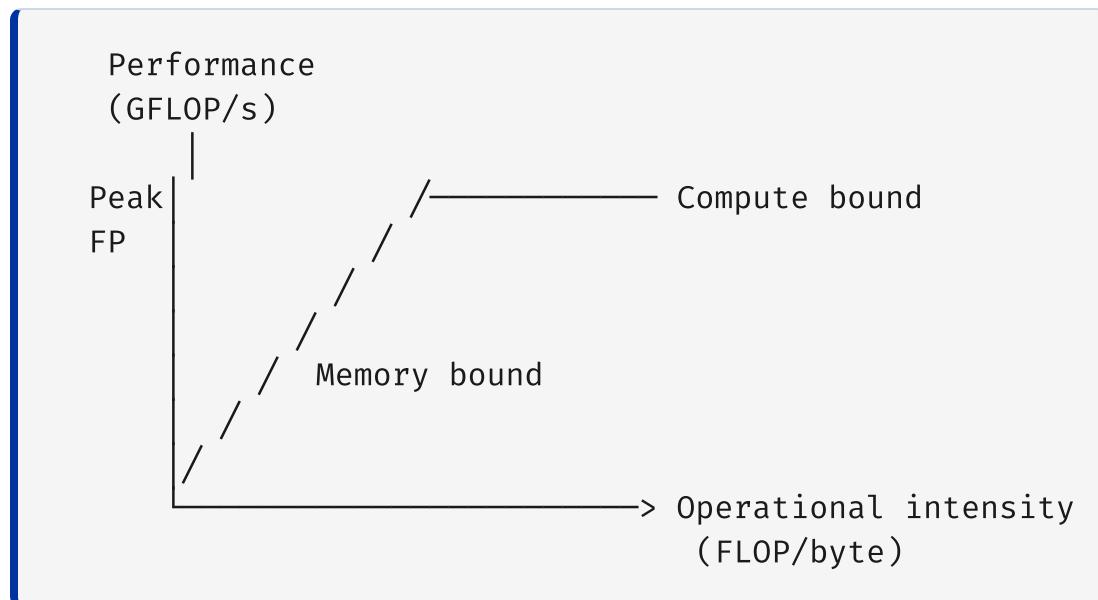
**Example:** Wide plateau at bottom = hot function consuming most time

# Roofline Model

---

## Are We Compute-Bound or Memory-Bound?

Roofline plot:



Interpretation:

- **Left of "roof":** Memory-bound → improve data layout
- **Right of "roof":** Compute-bound → algorithmic optimisation

# Measuring Operational Intensity

---

## FLOP/byte Ratio

```
# Use perf to count operations
$ perf stat -e fp_arith_inst_retired.scalar_single,\\
            fp_arith_inst_retired.128b_packed_double,\\
            cache-references,cache-misses \\
            ./program

# Calculate:
# FLOP/byte = (Total FP operations) / (Bytes accessed from memory)
```

### HEP typical values:

- Track fitting: ~10–50 FLOP/byte (compute-bound)
- Histogram filling: ~0.1–1 FLOP/byte (memory-bound)
- ROOT I/O: ~0.01 FLOP/byte (I/O-bound)

# Case Study: Optimising a Physics Kernel

---

## Starting Point

**Code:** Track parameter calculation

```
// Original implementation
void calculateParameters(const std::vector<Hit>& hits, TrackParams& params) {
    for (const auto& hit : hits) {
        Matrix3x3 cov = hit.covariance();
        Vector3 pos = hit.position();

        // Matrix operations ...
        Matrix3x3 invCov = cov.inverse();
        params.accumulate(pos, invCov);
    }
}
```

**Profile shows:** 60% of time in this function

# Case Study: Step 1 – Measure

---

```
$ perf record -g ./tracking benchmark.root  
$ perf report
```

```
Overhead  Symbol  
 60.2%  calculateParameters  
 15.3%  Matrix3x3 :: inverse  
 12.1%  Hit :: covariance  
  8.7%  Vector3 :: operator*
```

## Analysis:

- Main hotspot confirmed: `calculateParameters`
- Secondary hotspot: `Matrix3x3 :: inverse` (called inside loop)
- `Hit :: covariance` is also expensive

# Case Study: Step 2 – Analyse

Digging deeper with `perf annotate`:

```
$ perf annotate calculateParameters

calculateParameters():
 45.23%: mov    (%rax), %xmm0      # Loading from memory (slow!)
 12.34%: call   Matrix3x3::inverse # Function call overhead
  8.91%: movsd  %xmm0, (%rcx)
  ...
...
```

Findings:

1. Many cache misses (hits not contiguous)
2. Matrix inverse called every iteration (expensive)
3. Function call overhead

# Case Study: Step 3 – Optimise

---

## Improvement 1: Precompute Invariants

```
void calculateParameters(const std::vector<Hit>& hits, TrackParams& params) {
    // Precompute inverse covariances once
    std::vector<Matrix3x3> invCovs;
    invCovs.reserve(hits.size());

    for (const auto& hit : hits) {
        invCovs.push_back(hit.covariance().inverse());
    }
    // Now accumulate (no more inversions)
    for (size_t i = 0; i < hits.size(); ++i) {
        params.accumulate(hits[i].position(), invCovs[i]);
    }
}
```

Speedup: 1.8× (reduced inversions)

# Case Study: Step 3 – Optimise

---

## Improvement 2: SoA Data Layout

```
// Before: AoS
struct Hit {
    Vector3 position;
    Matrix3×3 covariance;
};
std::vector<Hit> hits;
// After: SoA
struct HitCollection {
    std::vector<float> x, y, z;
    std::vector<Matrix3×3> covariance;
};
```

**Result:** Better cache utilisation, vectorisation possible

**Speedup:** Additional 1.4x → **Total: 2.5x** vs original

# Case Study: Step 4 – Verify

---

```
# Before optimisation
$ perf stat ./tracking benchmark.root
Performance counter stats:
 10,234,567,890  cycles
      5,123,456,789  instructions      # 0.50 insns per cycle
      234,567,890  cache-misses
      5.234 seconds time elapsed

# After optimisation
$ perf stat ./tracking benchmark.root
Performance counter stats:
    4,123,456,789  cycles          # 2.5× faster
    6,234,567,890  instructions      # 1.51 insns per cycle
    98,765,432  cache-misses        # 2.4× fewer misses
    2.089 seconds time elapsed
```

# Optimisation Priorities

---

## Focus Your Efforts

### High Impact:

1. Algorithm choice
2. Data layout (AoS → SoA)
3. Removing allocations
4. Cache optimisation
5. Parallelisation

### Low Impact:

1. Variable naming
2. Inline hints (compiler knows)
3. Micro-optimisations
4. Premature vectorisation
5. Obscure tricks

**Rule:** 80/20 principle applies. First 20% effort gives 80% speedup.

# Common Optimisation Techniques (1/2)

---

## Quick Wins

### 1. Reserve vector capacity

```
std::vector<Track> tracks;  
tracks.reserve(estimatedSize); // Avoid reallocations
```

### 2. Pass by const reference

```
void process(const LargeObject& obj); // Not by value
```

# Common Optimisation Techniques (2/2)

---

## Quick Wins

### 3. Use `emplace_back` not `push_back`

```
tracks.emplace_back(args ...); // Construct in-place
```

### 4. Avoid unnecessary copies

```
auto result = compute(); // Move semantics
```

# Advanced Optimisation: Loop Unrolling

## Help the Compiler Vectorise

```
// Before (compiler may not unroll)
for (size_t i = 0; i < n; ++i) { result[i] = a[i] * b[i] + c[i]; }
// After (manual unroll by 4)
size_t i = 0;
for (; i + 4 ≤ n; i += 4) {
    result[i+0] = a[i+0] * b[i+0] + c[i+0];
    result[i+1] = a[i+1] * b[i+1] + c[i+1];
    result[i+2] = a[i+2] * b[i+2] + c[i+2];
    result[i+3] = a[i+3] * b[i+3] + c[i+3];
}
// Handle remainder
for (; i < n; ++i) { result[i] = a[i] * b[i] + c[i]; }
```

**Note:** Often compiler does this automatically with -O3

# When to Stop Optimising

---

Stop when:

1. **"Good enough" performance achieved** – Meets science requirements
2. **Diminishing returns** – Hours of work for 1% gain
3. **Code becoming unmaintainable** – Complexity not worth it
4. **Hardware bottleneck hit** – Memory bandwidth, disk I/O

**Remember:** Correct, maintainable code beats fast, broken code **every time**

# The Science Behind the Speed

---

## Higgs Discovery Impact

ATLAS 2012:

- Data to analyse:  $10 \text{ fb}^{-1}$
- Without optimisation: would have taken far longer
- With optimisation: months instead of years
- **Result: Announced July 4, 2012**

Optimised software enabled faster science.

# The Challenge Grows

---

Run 3 (2022–2026):

- Data rate: higher than Run 1
- CPU budget: limited growth
- **Need: major efficiency improvements**
- Achieved with techniques you're learning now

Your optimisations enable new discoveries.

10% speedup = months saved = faster science

# Optimisation Checklist

---

## Before You Start

- [ ] Profile to confirm hotspot (not guessing)
- [ ] Establish baseline measurements
- [ ] Write regression tests
- [ ] Document current performance
- [ ] Set a realistic target

## During Optimisation

- [ ] Change one thing at a time
- [ ] Profile after each change
- [ ] Verify correctness with tests
- [ ] Document what you tried (even failures)

# Summary: SD L3

---

## Key Takeaways:

1. **Workflow** – Measure → Analyse → Optimise → Verify
2. **Tools** – perf for profiling, flame graphs for visualisation
3. **Roofline model** – Identify if compute or memory bound
4. **Case study** – Real example of 2.5× speedup
5. **Priorities** – Algorithm and data layout first

**Next:** Hands-on exercises to apply profiling skills

# SOFTWARE DESIGN

---

## Exercises: E1 + E2

"Parallel programming: easier to talk about than to do"

2 hours total

Choose your own adventure

Parallelise event processing with OpenMP

GPU programming with CUDA

CPU vectorisation & cache optimisation

# Exercise SD E1: Parallelise a CPU Codepath

---

## Threading with OpenMP/TBB (60 minutes)

**Starting point:** Serial event processing loop

**Deliverable:** Scaling plot and bottleneck analysis

**Your tasks:**

1. Parallelise with OpenMP
2. Implement a reduction for counting
3. Measure scaling (1, 2, 4, 8 threads)
4. Diagnose the first bottleneck:
  - False sharing?
  - Load imbalance?
  - Grain size?

# SD E1: Getting Started

```
// File: event_processor.cpp (provided)

// Current serial implementation
void processEvents(const std::vector<Event>& events) {
    int totalTracks = 0;

    for (const auto& event : events) {
        auto tracks = event.reconstructTracks();
        totalTracks += tracks.size();

        // More processing ...
    }

    std::cout << "Total tracks: " << totalTracks << std::endl;
}
```

**Your job:** Parallelise this and measure speedup

# SD E1: Hints (1/2)

---

## Parallel Implementation

```
void processEvents(const std::vector<Event>& events) {
    int totalTracks = 0;

    #pragma omp parallel for reduction(+:totalTracks)
    for (size_t i = 0; i < events.size(); ++i) {
        auto tracks = events[i].reconstructTracks();
        totalTracks += tracks.size();
    }

    std::cout << "Total tracks: " << totalTracks << std::endl;
}
```

## SD E1: Hints (2/2)

---

### Measure scaling

```
for threads in 1 2 4 8; do  
    OMP_NUM_THREADS=$threads ./event_processor  
done
```

# SD E1: What to Look For

---

Ideal scaling:

```
1 thread: 10.0 seconds (baseline)
2 threads: 5.0 seconds (2.0× speedup)
4 threads: 2.5 seconds (4.0× speedup)
8 threads: 1.25 seconds (8.0× speedup)
```

Reality (you'll see something like):

```
1 thread: 10.0 seconds
2 threads: 5.8 seconds (1.7× speedup) ← Why not 2×?
4 threads: 3.2 seconds (3.1× speedup) ← Why not 4×?
8 threads: 2.1 seconds (4.8× speedup) ← Why not 8×?
```

Your analysis should identify: False sharing? Memory bandwidth? Load imbalance?

# Exercise SD E2: Advanced Path

---

## Choose One (60 minutes)

### Option A: GPU Porting (if GPUs are available)

- Port hot loop to CUDA/HIP
- Manage host-device transfers
- Perform occupancy analysis
- Compare CPU vs GPU performance

### Option B: CPU Vectorisation & Tiling (CPU-only)

- Vectorise critical loop
- Apply tiling for cache optimisation
- Use the roofline model to understand limits
- Show before/after perf counters

# SD E2 Option A: GPU Basics – CUDA Example

```
// CPU version
void add(float* a, float* b, float* c, int n) {
    for (int i = 0; i < n; ++i) {
        c[i] = a[i] + b[i];
    }
}

// GPU version
__global__ void add_kernel(float* a, float* b, float* c, int n) {
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    if (i < n) {
        c[i] = a[i] + b[i];
    }
}

// Host code
add_kernel<<<(n+255)/256, 256>>>(d_a, d_b, d_c, n);
```

# SD E2 Option B: Vectorisation | Manual SIMD

```
#include <immintrin.h> // AVX2

// Vectorised addition (8 floats at once)
void add_avx(float* a, float* b, float* c, size_t n) {
    size_t i = 0;

    // Process 8 floats at a time
    for (; i + 8 ≤ n; i += 8) {
        __m256 va = _mm256_loadu_ps(&a[i]);
        __m256 vb = _mm256_loadu_ps(&b[i]);
        __m256 vc = _mm256_add_ps(va, vb);
        _mm256_storeu_ps(&c[i], vc);
    }
    // Handle remainder
    for (; i < n; ++i) {
        c[i] = a[i] + b[i];
    }
}
```

# Both Options: Deliverables

---

## What to Submit

### Documentation:

- [ ] Performance measurements (before/after)
- [ ] Scaling analysis
- [ ] Bottleneck identification
- [ ] Lessons learned

### Code:

- [ ] Clean, commented implementation
- [ ] Tests verifying correctness
- [ ] Build system integration
- [ ] CI running benchmarks

### Bonus:

- [ ] Roofline analysis (Option B)
- [ ] Occupancy analysis (Option A)
- [ ] Comparison with theoretical peak

# CROSS-CUTTING: INFRASTRUCTURE

---

The WLCG Stack

15 minutes of context

# The WLCG Stack

---

## Worldwide LHC Computing Grid

**Problem:** LHC data is too big for one site

**Solution:** Distributed computing across ~170 sites worldwide

**Key components:**

1. **EOS** – High-throughput storage
2. **CVMFS** – Software distribution
3. **Apptainer** – Container runtime
4. **HTCondor/Slurm** – Job scheduling

# EOS: CERN Disk Storage

---

## Exabyte-Scale File System

### Characteristics:

- High latency (~10-100ms per operation)
- High throughput (GB/s aggregate)
- Designed for large sequential I/O

**Use case:** Store physics data files (ROOT, HDF5)

### Why we buffer I/O:

```
// Bad: Many small reads (slow on EOS)
for (int i = 0; i < n; ++i) {
    float value;
    file.read(&value, sizeof(float));
}

// Good: One large read (fast on EOS)
std::vector<float> buffer(n);
file.read(buffer.data(), n * sizeof(float));
```

# CVMFS: Software Distribution

---

## CernVM File System

**Purpose:** Distribute software to all grid sites without copying

### How it works:

- Read-only file system
- HTTP-based content delivery
- Aggressive caching (via Squid)
- Cryptographic integrity

### Mounted at:

- /cvmfs/sft.cern.ch/
- /cvmfs/atlas.cern.ch/
- etc.

### Example: Access LCG software releases

```
source /cvmfs/sft.cern.ch/lcg/views/LCG_104/x86_64-el9-gcc13-opt/setup.sh  
root -l # Now have ROOT, Geant4, etc. available
```

# Apptainer (formerly Singularity)

---

## Container Runtime for HPC

### Why not Docker?

- Docker requires root privileges
- Not allowed on shared grid sites

**Apptainer:** Rootless container runtime

```
# Pull container
apptainer pull docker://rootproject/root:latest

# Run analysis inside container
apptainer exec root_latest.sif root -l analysis.C

# Mount /cvmfs inside container
apptainer exec --bind /cvmfs root_latest.sif \
    source /cvmfs/sft.cern.ch/lcg/views/LCG_104/setup.sh
```

# Why We Need Relocatable Binaries

---

## The Distribution Challenge

**Problem:** Software built on one site might not run on another

**Causes:**

- Different library paths (`/usr/lib` vs `/opt/lib`)
- Different compiler versions
- Different OS versions

**Solutions:**

1. **CVMFS:** All sites mount same software
2. **Containers:** Bundle dependencies
3. **Relocatable builds:** Software works regardless of install path

# Putting It Together

---

## Typical HEP Workflow on Grid

```
# 1. Submit job to HTCondor  
condor_submit job.sub  
  
# 2. Job runs on a random grid site  
# 3. Worker node:  
#     - Mounts CVMFS for software  
#     - Pulls container if needed  
#     - Reads data from EOS via XRootD  
#     - Processes events  
#     - Writes output back to EOS  
  
# 4. Output available globally  
xrdcp root://eosatlas.cern.ch//path/to/output.root ./
```

**Key point:** Same job runs anywhere on the grid thanks to CVMFS + containers

# CONCLUSION

---

Bringing It All Together

# What We've Covered

---

## Tools & Techniques

- Correctness with testing and sanitisers
- Collaboration with Git and CI
- Performance fundamentals

## Software Design

- Many-core programming
- Parallel patterns and data layout
- Profiling and optimisation

## Infrastructure

- WLCG stack overview
- CVMFS, EOS, Apptainer



# Key Principles to Remember (1/2)

## 1. Correctness First

- Test everything
- Use sanitisers
- Don't sacrifice correctness for speed

## 2. Measure Before Optimising

- Profile to find real hotspots
- Don't guess
- Verify every change

## 3. Collaborate Effectively

- Atomic commits
- CI as production
- Code review is crucial
- Document decisions
- Ask for help & share knowledge

# Key Principles to Remember (2/2)

---

## 4. Think About Data Layout

- SoA usually beats AoS
- Cache locality matters
- Pad to avoid false sharing

## 5. Use the Right Tool

- OpenMP for simple parallelism
- TBB for complex task graphs
- Profile with perf
- Leverage HEP frameworks (ROOT, CMSSW)

# Beyond this course & this School

---

## Continue Learning

### Resources:

- Books:
  - "Computer Systems: A Programmer's Perspective" (Bryant & O'Hallaron)
  - "The Art of Multiprocessor Programming" (Herlihy & Shavit)
- Online:
  - CERN/HSF Software Training modules
  - C++ Core Guidelines
  - Intel's optimisation manuals
- Practice:
  - Contribute to CMSSW, ROOT, Geant4
  - Optimise your own analysis code
  - Profile real workloads

# Final Thoughts

---

Writing high-performance code for HEP is both:

- **Engineering challenge** – Technical skills matter
- **Scientific discipline** – Correctness is paramount

Hopefully I did a decent job and you now have the tools to:

- Write correct, tested code
- Collaborate effectively in large teams
- Profile and optimise systematically
- Leverage modern parallel hardware

Most importantly: You understand **why** these practices matter for science.

# Thank You!

---

You're Part of History

47th CSC, 1st in Latin America!

You now can:

- Write correct, tested code
- Collaborate in global teams
- Profile and optimise
- Use parallel hardware
- Work with real HEP data

You understand: These skills enable better science.

# The Bigger Picture

---

Remember:

- CERN Open Data makes knowledge accessible
- Same tools discovered the Higgs boson
- Your contributions impact researchers worldwide
- Computing skills enable scientific discovery

"International cooperation is possible and necessary for equity and progress in education."

# Your Impact Starts Now

---

## Making a Difference

"International cooperation is possible and necessary for equity and progress in education around the world."

### What you can do next:

- Contribute to ATLAS Open Data
- Join HSF working groups
- Mentor students in your country
- Apply these skills in your research
- Share what you learned

### The HEP community:

- Is **global** (40+ countries)
- Is **collaborative** (sharing > competing)
- Is **always learning** (~2 weeks at schools like this!)
- **Needs you** to make it better

**Let's stay connected.** The best discoveries happen when we work together.

# What's Next?

---

**Week 1:** Set up dev container, add tests, enable sanitisers, profile something

**Month 1:** Implement CI, parallelise a loop, and contribute a bug fix

**This Year:** Attend HSF training, present your work, mentor someone

# Questions & Contact

---

During school:

- Ask anytime
- Help each other
- Share what you learn

After:

- <https://www.linkedin.com/in/arturo-sanchez-pineda/>
- <https://github.com/artfisica/csc2026/>

# Join the Community

---

## Resources:

- HSF Training: <https://hepsoftwarefoundation.org/training/>
- ATLAS Open Data: <https://opendata.atlas.cern>
- CERN Docs: <https://docs.cern.ch/>
- GitHub: <https://github.com/HSF>
- CERN Mattermost: Join CERN community channels if you can :)

Remember: every expert was once a beginner.

# Quick Reference

---

## Testing & Debugging:

```
cmake -DCMAKE_CXX_FLAGS="-fsanitize=address" ..  
ctest --output-on-failure
```

## Git:

```
git status  
git add -p  
git commit -m "Fix: ... "
```

## Profiling:

```
perf record -g ./program  
perf report
```

# BACKUP SLIDES

---

Additional Reference Material

# Backup: CMake Basics

```
cmake_minimum_required(VERSION 3.16)
project(HEPAnalysis CXX)
set(CMAKE_CXX_STANDARD 17)

find_package(ROOT REQUIRED COMPONENTS Core RIO Tree)
find_package(Catch2 REQUIRED)

add_library(analysis SHARED src/Particle.cpp)
target_link_libraries(analysis PUBLIC ROOT::Core ROOT::Tree)

add_executable(analyze src/main.cpp)
target_link_libraries(analyze analysis)

add_executable(tests tests/test_particle.cpp)
target_link_libraries(tests analysis Catch2::Catch2WithMain)
enable_testing()
add_test(NAME unit_tests COMMAND tests)
```

# Backup: GDB Cheat Sheet

---

Command	Action
<code>run</code>	Start execution
<code>break main</code>	Set breakpoint
<code>continue</code>	Resume after breakpoint
<code>next / step</code>	Step over / into
<code>print var</code>	Print variable
<code>backtrace</code>	Show call stack
<code>frame N</code>	Switch to frame N
<code>watch var</code>	Break when var changes

Start: `gdb ./program` or `gdb --args ./program arg1`

# Backup: Compiler Flags

---

Flag	Purpose
-O0	No optimisation (debug)
-O2	Good balance
-O3	Aggressive
-march=native	Optimise for this CPU
-g	Debug symbols
-Wall -Wextra	Enable warnings
-fsanitize=address	ASan

Recommended: -O2 -g -Wall -Wextra -march=native

# Backup: ROOT Macro Tips

```
void analysis() {
    // Modern way: RDataFrame
    ROOT::RDataFrame df("Events", "data.root");
    auto h = df.Filter("nJets ≥ 4")
        .Histo1D({"h", "pT", 100, 0, 200}, "pt");
    h→Draw();

    // Save output
    auto out = TFile::Open("output.root", "RECREATE");
    h→Write();
    out→Close();
}
```

Run: root -l analysis.C

# Backup: Python Performance

```
import numpy as np

# ❌ Slow: Python loops
result = [x**2 for x in data]

# ✅ Fast: NumPy vectorisation
result = np.array(data) ** 2    # 100× faster!

# ✅ Generators for large data
def process(filename):
    with open(filename) as f:
        for line in f:
            yield transform(line)
```

Profile: `python -m cProfile script.py`

# Backup: Performance At-a-Glance

## Optimisation Priority:

Impact	Speedup
Algorithm	100–1000×
Data Structure	10–100×
Memory Layout	5–50×
Parallelisation	4–8×
Compiler -O3	2–5×
Micro-opts	1.001×

## Memory Hierarchy:

Level	Cycles	Size
Registers	1	1 KB
L1 Cache	4	64 KB
L2 Cache	12	1 MB
L3 Cache	40	32 MB
RAM	200	128 GB
SSD	50000	2 TB

Start from the top!

RAM is 200× slower than L1!

# Backup: Further Reading

---

## Books:

- "Effective Modern C++" — Scott Meyers
- "C++ Concurrency in Action" — Anthony Williams
- "Systems Performance" — Brendan Gregg

**Communities:** CERN Mattermost, HSF Forum

## Further Reading – Online

---

- ROOT documentation – <https://root.cern/>
- HSF Training – <https://hepsoftwarefoundation.org/training/>
- C++ Core Guidelines – <https://isocpp.github.io/CppCoreGuidelines/>

## REFERENCES – Core HEP & WLCG

---

- Worldwide LHC Computing Grid (WLCG): <https://home.cern/science/computing/grid>
- CERN Open Data Portal: <https://opendata.cern.ch/>
- ATLAS Open Data: <https://opendata.atlas.cern/>
- ROOT Framework (home): <https://root.cern/>

# REFERENCES – Testing & Debugging

---

- Catch2 (C++ testing): <https://github.com/catchorg/Catch2>
- pytest (Python testing): <https://docs.pytest.org/en/stable/>
- AddressSanitizer (ASan): <https://clang.llvm.org/docs/AddressSanitizer.html>
- UndefinedBehaviorSanitizer (UBSan):  
<https://clang.llvm.org/docs/UndefinedBehaviorSanitizer.html>
- ThreadSanitizer (TSan): <https://clang.llvm.org/docs/ThreadSanitizer.html>
- clang-tidy: <https://clang.llvm.org/extras/clang-tidy/>
- C++ Core Guidelines: <https://isocpp.github.io/CppCoreGuidelines/CppCoreGuidelines>

# REFERENCES – Performance & Profiling

---

- Google Benchmark: <https://github.com/google/benchmark>
- Linux perf (kernel docs): <https://www.kernel.org/doc/html/latest/admin-guide/perf.html>
- FlameGraph: <https://github.com/brendangregg/FlameGraph>
- Roofline Model (LBL): <https://crd.lbl.gov/departments/computer-science/PAR/research/roofline/>

# REFERENCES – Parallelism & Frameworks

---

- OpenMP (official): <https://www.openmp.org/>
- oneTBB (docs): <https://oneapi-src.github.io/oneTBB/>
- ROOT RDataFrame (docs): [https://root.cern/doc/master/group\\_\\_dataframe.html](https://root.cern/doc/master/group__dataframe.html)
- EnableImplicitMT (ROOT): <https://root.cern/doc/master/namespacEROOT.html>
- CMSSW (CMS Software): <https://cms-sw.github.io/>
- Geant4 – Multi-threading: <https://geant4-userdoc.web.cern.ch/UsersGuides/ForApplicationDeveloper/html/GettingStarted/multithreading.html>

# REFERENCES – Infrastructure

---

- CernVM-FS (CVMFS): <https://cernvm.cern.ch/fs/>
- EOS (CERN storage): <https://eos-docs.web.cern.ch/>
- Apptainer (Singularity): <https://apptainer.org/docs/>
- XRootD: <https://xrootd.slac.stanford.edu/>

# REFERENCES – Training & Community

---

- HSF Training Center: <https://hepsoftwarefoundation.org/training/>
- ROOT User Guides: <https://root.cern/doc/master/index.html>
- CERN Docs: <https://docs.cern.ch/>
- Communities: CERN Mattermost, HSF Forum

# Thank You!

---

CSC Latin America 2026 – Santiago, Chile

Arturo Sánchez Pineda – INAIT AI | Creative Commons Venezuela